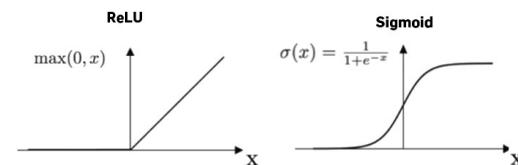
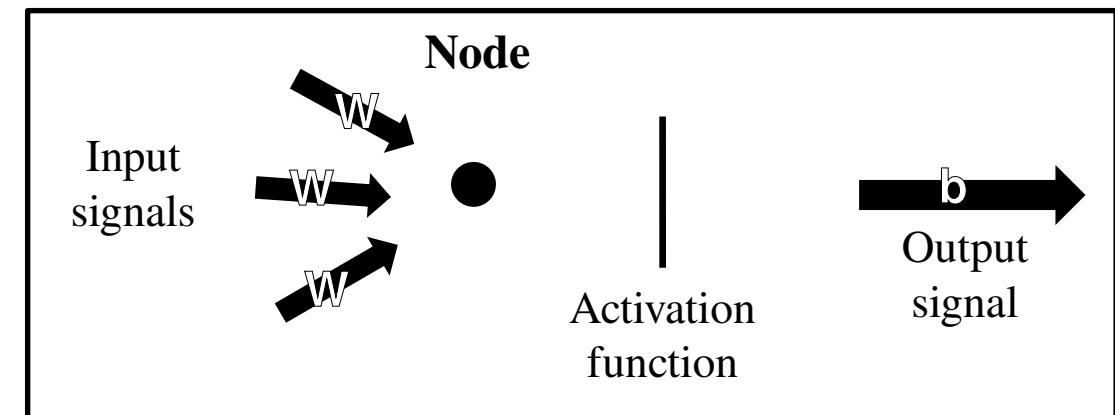
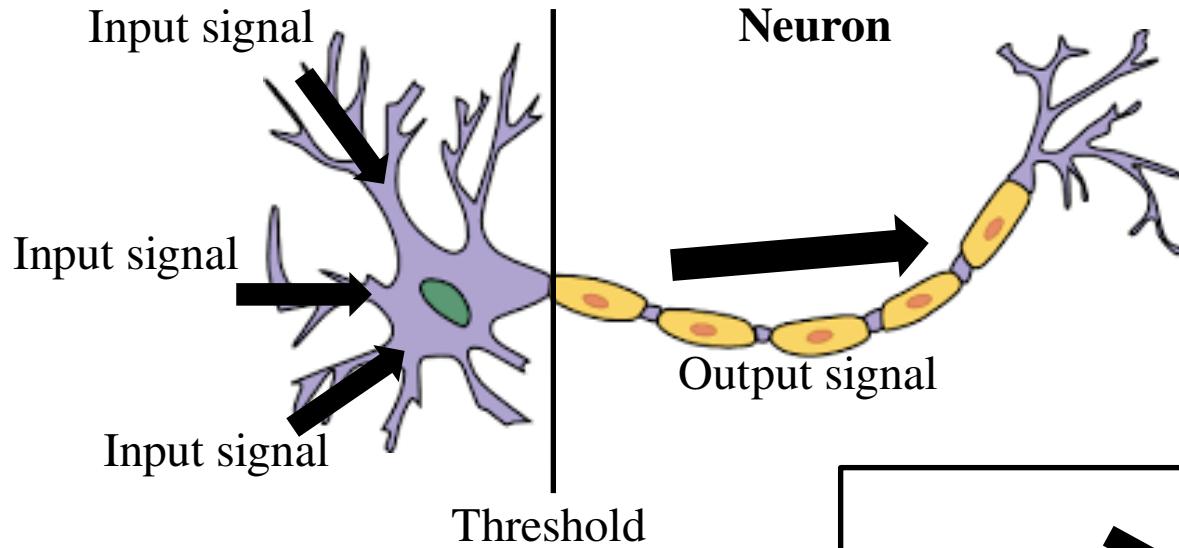


Using deep neural networks as a modelling framework to understand ventral stream computations

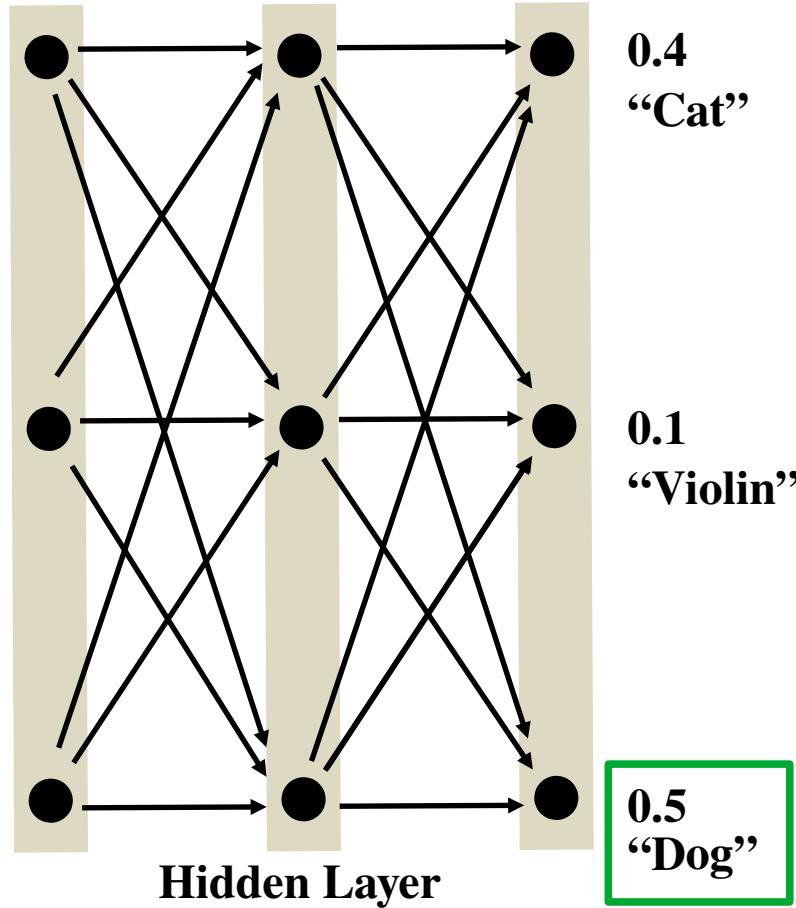
Emer Jones

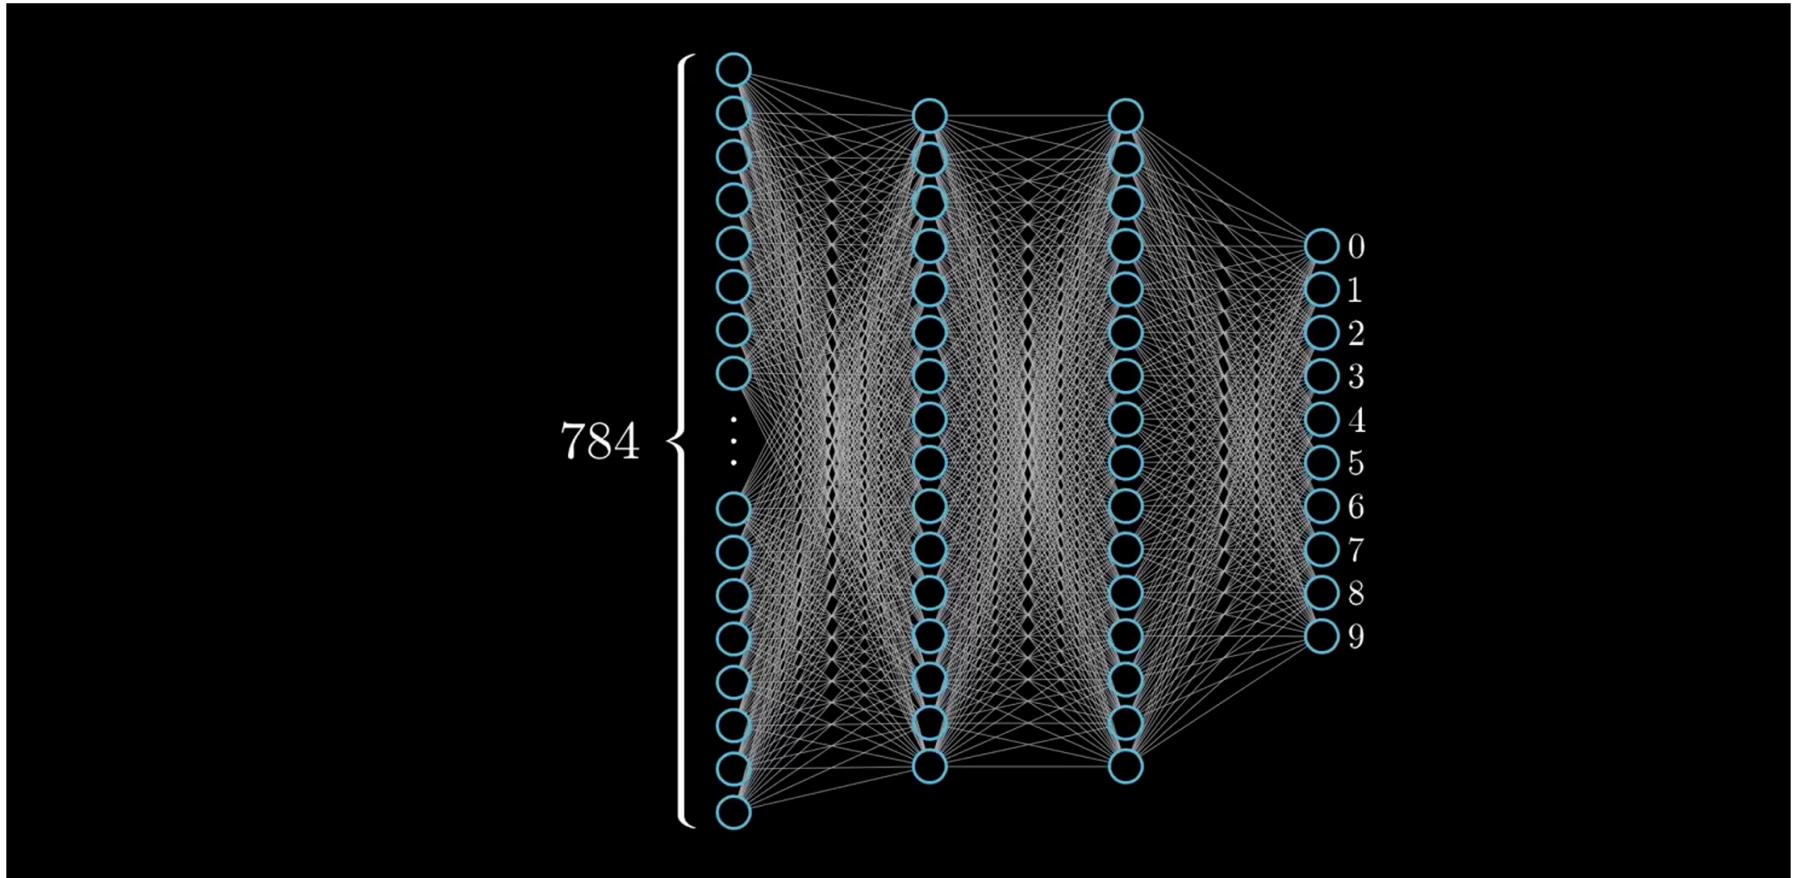
PhD student
MRC Cognition and Brain Sciences Unit,
University of Cambridge

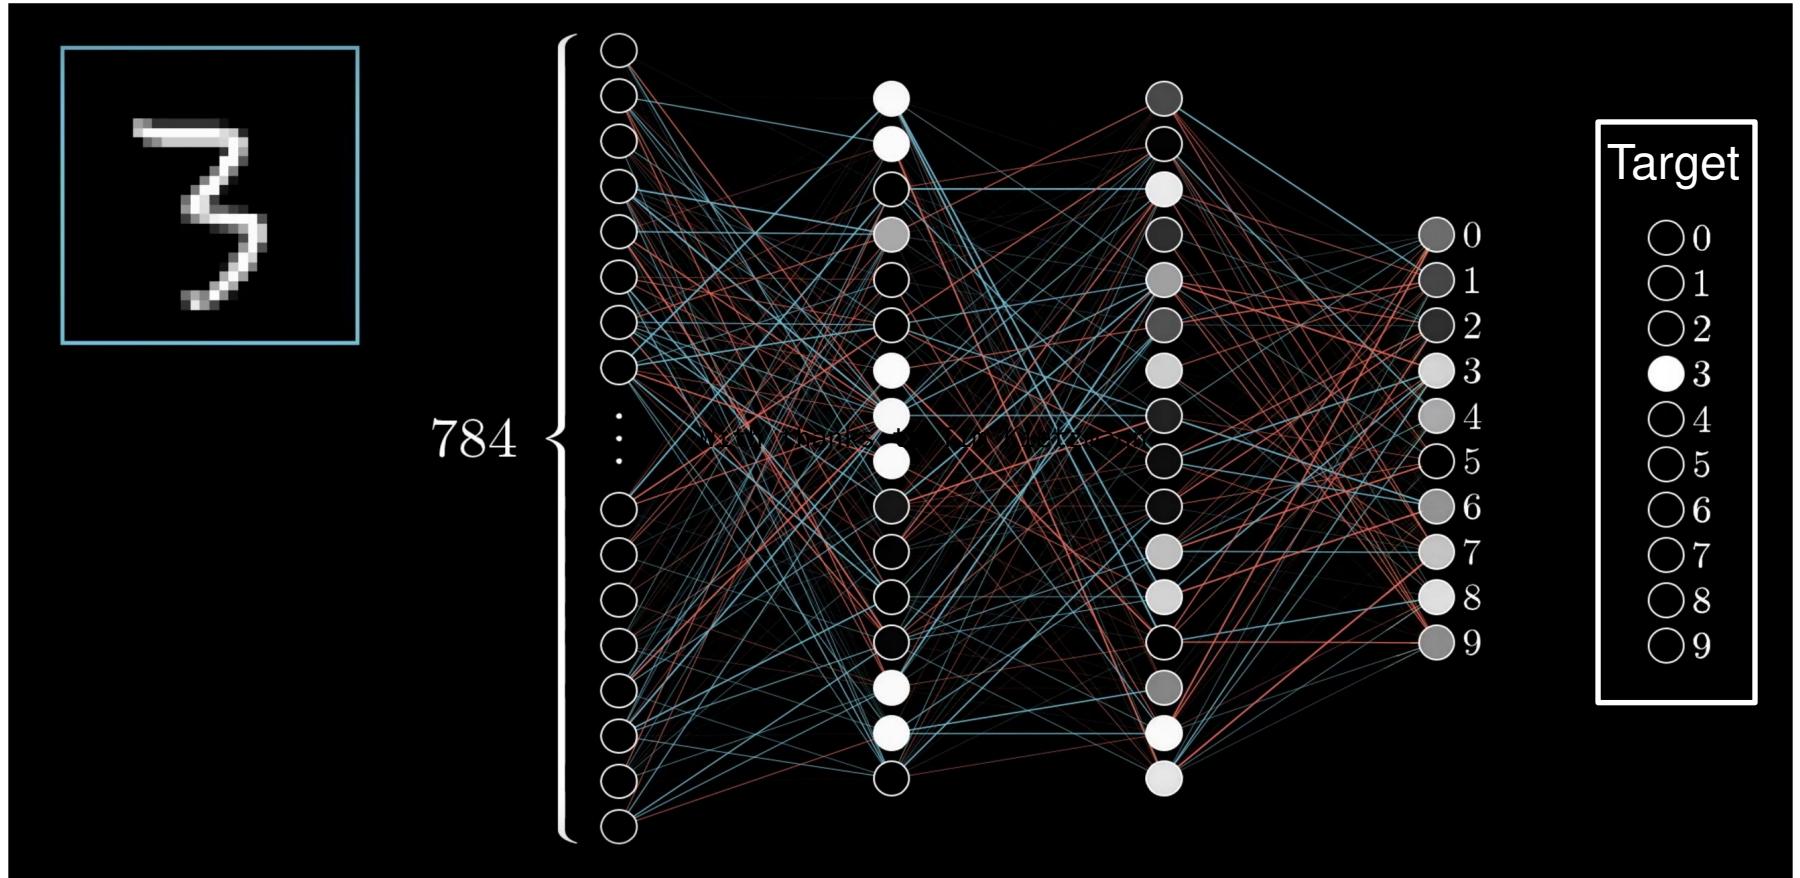




Input Layer:
Pixels of an image







```
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Dense

# Load MNIST data and reshape from 3D to 4D
(x_train, y_train), (x_test, y_test) = mnist.load_data()

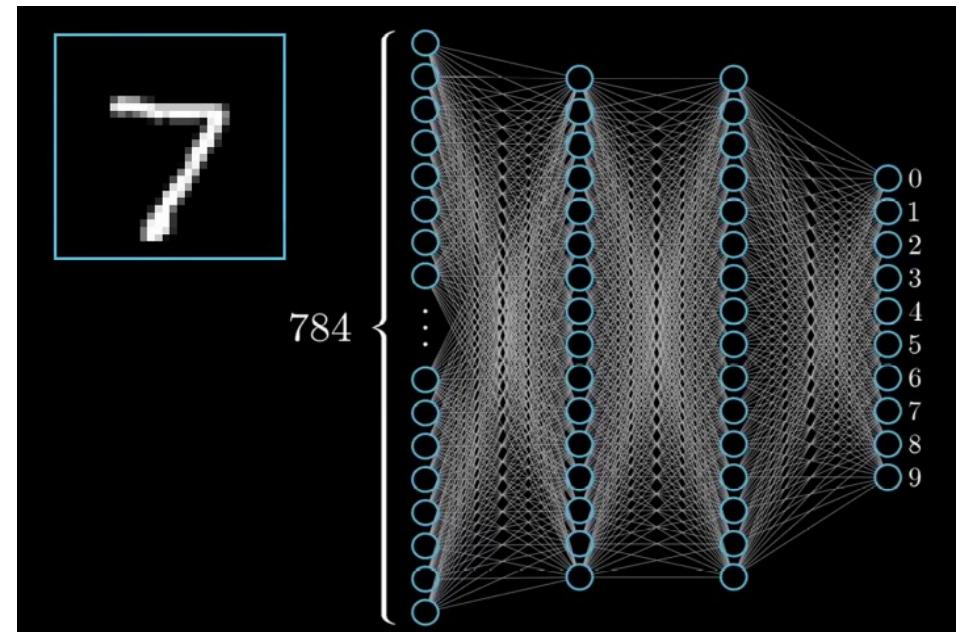
x_train = x_train.reshape(x_train.shape[0], 784)
x_test = x_test.reshape(x_test.shape[0], 784)

# Define the model architecture
inputs = Input(shape=(784,))
x = Dense(16)(inputs)
x = Dense(16)(x)
readout = Dense(10, activation=tf.nn.softmax)(x)
model = Model(inputs=inputs, outputs=readout)

# Define which loss & optimisation functions and metric(s) to use
model.compile(loss='sparse_categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

# Train and evaluate the model
model.fit(x=x_train, y=y_train, epochs=10)
model.evaluate(x_test, y_test)
```

x_train: (60000, 784)
y_train: (60000)
x_test: (10000, 784)
y_test: (10000)



Number of parameters: 13,002
Accuracy: 87%

```
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Dense

# Load MNIST data and reshape from 3D to 4D
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(x_train.shape[0], 784)
x_test = x_test.reshape(x_test.shape[0], 784)

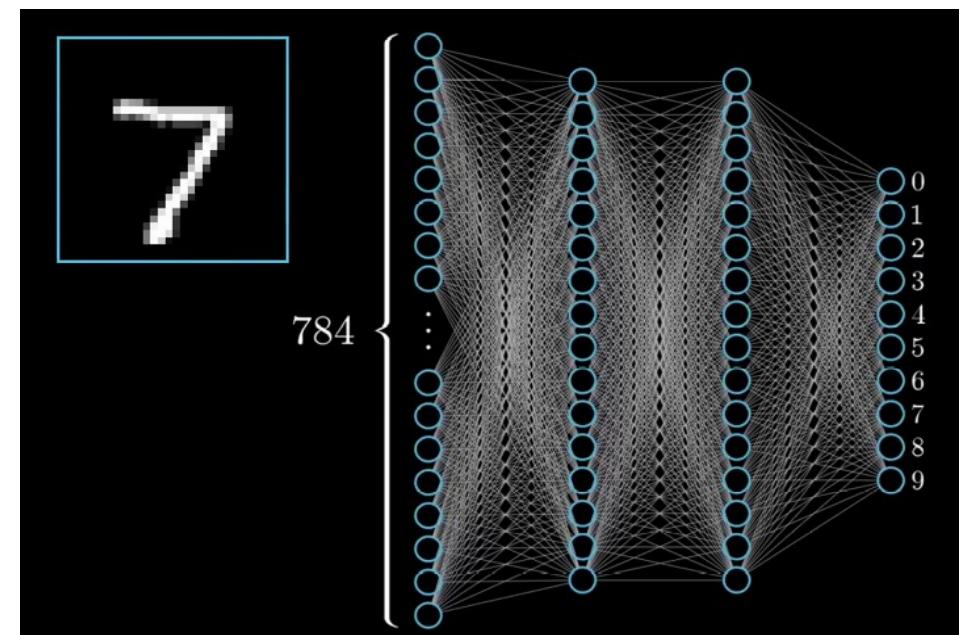
# Rescale pixel values between 0 and 255 to values between 0 and 1
x_train = x_train.astype('float32')/255
x_test = x_test.astype('float32')/255

# Define the model architecture
inputs = Input(shape=(784,))
x = Dense(16)(inputs)
x = Dense(16)(x)
readout = Dense(10, activation=tf.nn.softmax)(x)
model = Model(inputs=inputs, outputs=readout)

# Define which loss & optimisation functions and metric(s) to use
model.compile(loss='sparse_categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

# Train and evaluate the model
model.fit(x=x_train, y=y_train, epochs=10)
model.evaluate(x_test, y_test)
```

x_train: (60000, 784)
y_train: (60000)
x_test: (10000, 784)
y_test: (10000)



Number of parameters: 13,002
Accuracy: 93%

```

import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Dense

# Load MNIST data and reshape from 3D to 4D
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(x_train.shape[0], 784)
x_test = x_test.reshape(x_test.shape[0], 784)

# Rescale pixel values between 0 and 255 to values between 0 and 1
x_train = x_train.astype('float32')/255
x_test = x_test.astype('float32')/255

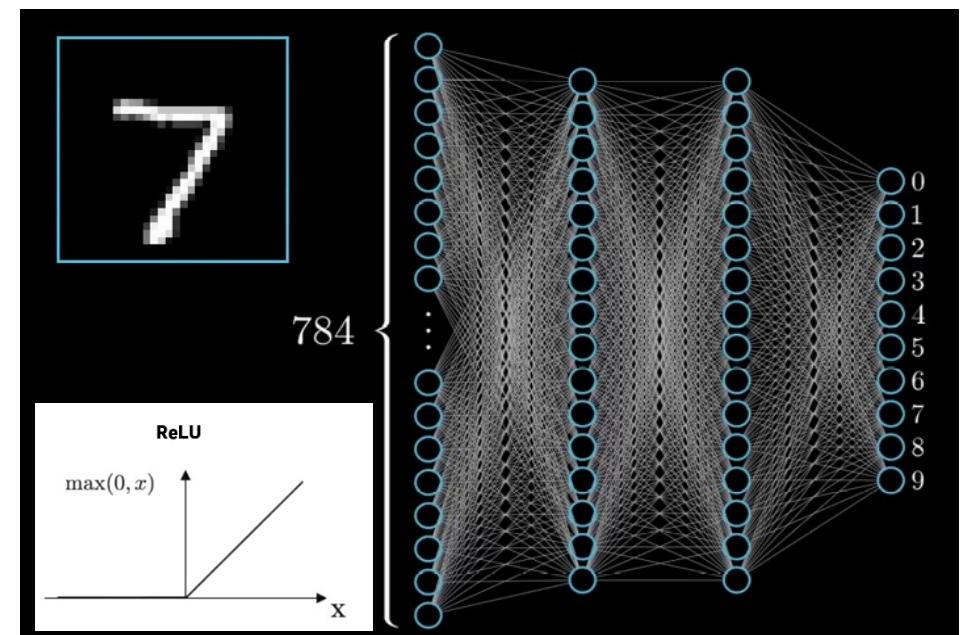
# Define the model architecture
inputs = Input(shape=(784,))
x = Dense(16, activation=tf.nn.relu)(inputs)
x = Dense(16, activation=tf.nn.relu)(x)
readout = Dense(10, activation=tf.nn.softmax)(x)
model = Model(inputs=inputs, outputs=readout)

# Define which loss & optimisation functions and metric(s) to use
model.compile(loss='sparse_categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

# Train and evaluate the model
model.fit(x=x_train, y=y_train, epochs=10)
model.evaluate(x_test, y_test)

```

x_train: (60000, 784)
y_train: (60000)
x_test: (10000, 784)
y_test: (10000)



Number of parameters: 13,002
Accuracy: 96%

```

import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Dense, Conv2D, Flatten, MaxPooling2D

# Load MNIST data and reshape from 3D to 4D
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(x_train.shape + (1,))
x_test = x_test.reshape(x_test.shape + (1,))

# Rescale pixel values between 0 and 255 to values between 0 and 1
x_train = x_train.astype('float32')/255
x_test = x_test.astype('float32')/255

# Define the model architecture
inputs = Input(shape=(28,28,1))
x = Conv2D(8, kernel_size=(3,3))(inputs)
x = MaxPooling2D(pool_size=(2,2))(x)
x = Conv2D(8, kernel_size=(3,3))(x)
x = MaxPooling2D(pool_size=(2,2))(x)
x = Flatten()(x)
readout = Dense(10, activation=tf.nn.softmax)(x)
model = Model(inputs=inputs, outputs=readout)

# Define which loss & optimisation functions and metric(s) to use
model.compile(loss='sparse_categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

# Train and evaluate the model
model.fit(x=x_train, y=y_train, epochs=10)
model.evaluate(x_test, y_test)

```

x_train: (60000, 28, 28, 1)

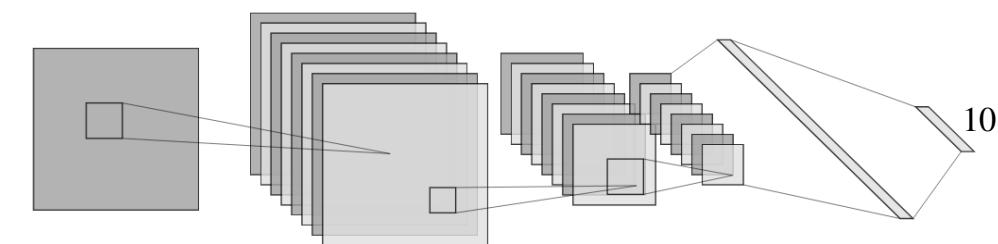
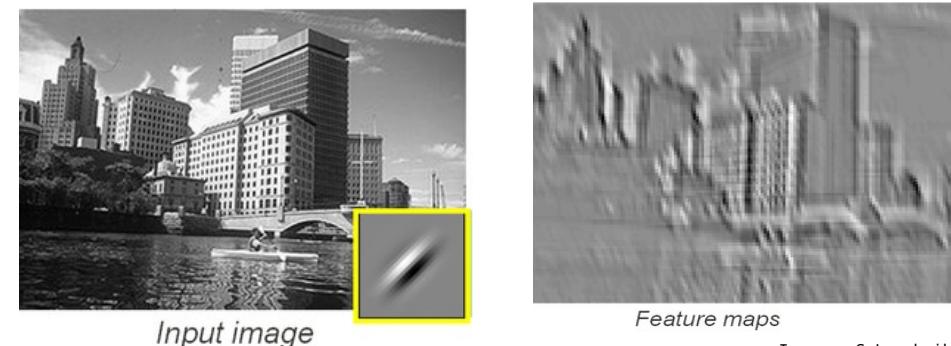


Diagram: <http://alexlenail.me/NN-SVG/LeNet.html>

Number of parameters: 2,674
Accuracy: 98%

```

import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Dense, Conv2D, Flatten, MaxPooling2D

# Load MNIST data and reshape from 3D to 4D
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(x_train.shape + (1,))
x_test = x_test.reshape(x_test.shape + (1,))

# Rescale pixel values between 0 and 255 to values between 0 and 1
x_train = x_train.astype('float32')/255
x_test = x_test.astype('float32')/255

# Define the model architecture
inputs = Input(shape=(28,28,1))
x = Conv2D(8, kernel_size=(3,3))(inputs)
x = MaxPooling2D(pool_size=(2,2))(x)
x = Conv2D(8, kernel_size=(3,3))(x)
x = MaxPooling2D(pool_size=(2,2))(x)
x = Flatten()(x)
readout = Dense(10, activation=tf.nn.softmax)(x)
model = Model(inputs=inputs, outputs=readout)
model.summary()

# Define which loss & optimisation functions and metric(s) to use
model.compile(loss='sparse_categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

# Train and evaluate the model
model.fit(x=x_train, y=y_train, epochs=10)
model.evaluate(x_test, y_test)

```

```

Train on 60000 samples
Epoch 1/10
60000/60000 [=====] - 9s 147us/sample - loss: 0.3278 - accuracy: 0.9045
Epoch 2/10
60000/60000 [=====] - 9s 147us/sample - loss: 0.1252 - accuracy: 0.9628
Epoch 3/10
60000/60000 [=====] - 9s 155us/sample - loss: 0.0967 - accuracy: 0.9707
Epoch 4/10
60000/60000 [=====] - 9s 150us/sample - loss: 0.0825 - accuracy: 0.9750
Epoch 5/10
60000/60000 [=====] - 10s 159us/sample - loss: 0.0731 - accuracy: 0.9781
Epoch 6/10
60000/60000 [=====] - 9s 150us/sample - loss: 0.0685 - accuracy: 0.9790
Epoch 7/10
60000/60000 [=====] - 9s 147us/sample - loss: 0.0631 - accuracy: 0.9809
Epoch 8/10
60000/60000 [=====] - 9s 147us/sample - loss: 0.0606 - accuracy: 0.9810
Epoch 9/10
60000/60000 [=====] - 9s 150us/sample - loss: 0.0567 - accuracy: 0.9827
Epoch 10/10
60000/60000 [=====] - 9s 144us/sample - loss: 0.0541 - accuracy: 0.9831

10000/10000 [=====] - 1s 71us/sample - loss: 0.0565 - accuracy: 0.9822

```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[None, 28, 28, 1]	0
conv2d (Conv2D)	(None, 26, 26, 8)	80
max_pooling2d (MaxPooling2D)	(None, 13, 13, 8)	0
conv2d_1 (Conv2D)	(None, 11, 11, 8)	584
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 8)	0
flatten (Flatten)	(None, 200)	0
dense (Dense)	(None, 10)	2010

Total params: 2,674
 Trainable params: 2,674
 Non-trainable params: 0

```
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Dense, Conv2D, Flatten, MaxPooling2D
from matplotlib import image

# Load MNIST data and reshape from 3D to 4D
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(x_train.shape + (1,))
x_test = x_test.reshape(x_test.shape + (1,))

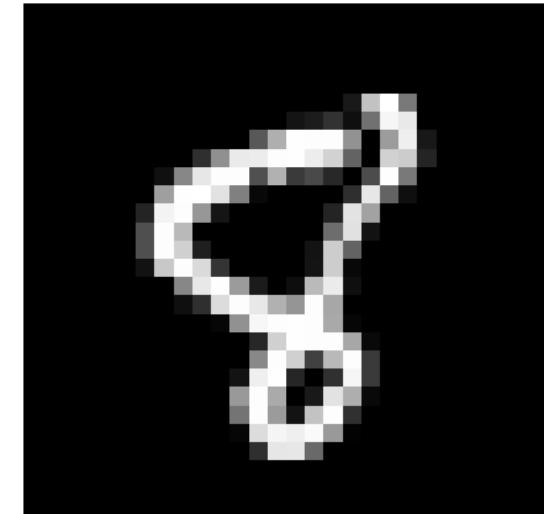
# Rescale pixel values between 0 and 255 to values between 0 and 1
x_train = x_train.astype('float32')/255
x_test = x_test.astype('float32')/255

# Define the model architecture
inputs = Input(shape=(28,28,1))
x = Conv2D(8, kernel_size=(3,3))(inputs)
x = MaxPooling2D(pool_size=(2,2))(x)
x = Conv2D(8, kernel_size=(3,3))(x)
x = MaxPooling2D(pool_size=(2,2))(x)
x = Flatten()(x)
readout = Dense(10, activation=tf.nn.softmax)(x)
model = Model(inputs=inputs, outputs=readout)
model.summary()

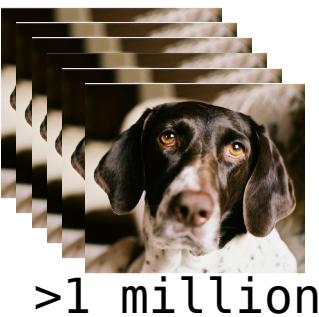
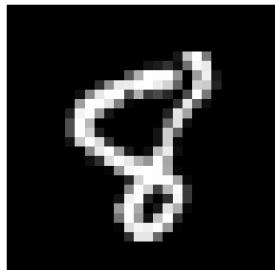
# Define which loss & optimisation functions and metric(s) to use
model.compile(loss='sparse_categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

# Train and evaluate the model
model.fit(x=x_train, y=y_train, epochs=10)
model.evaluate(x_test, y_test)

# Single image inference
test_image = image.imread('test_img.png')
pred = model.predict(test_image.reshape(1, 28, 28, 1))
print(pred.argmax())
```



Predicted digit for test image: 8



>1 million

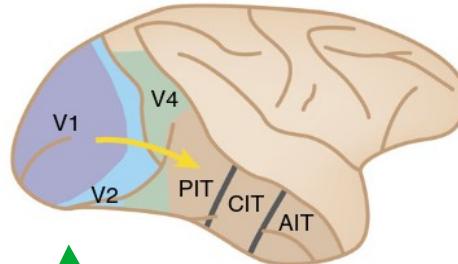


10 classes

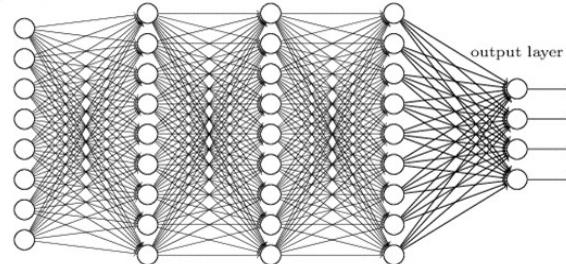


100s/1000s
of classes

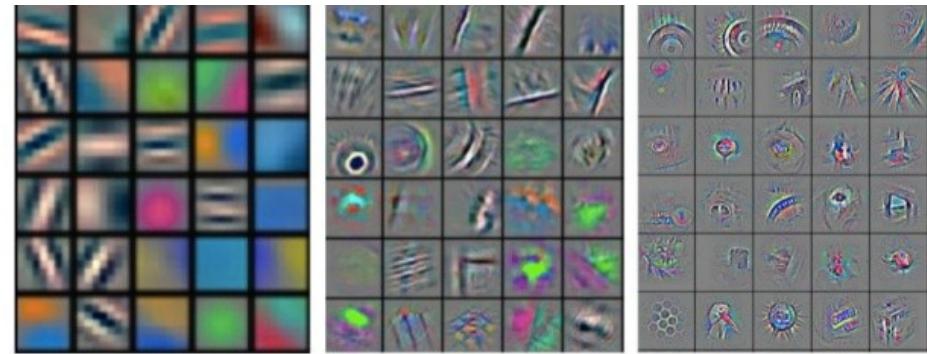
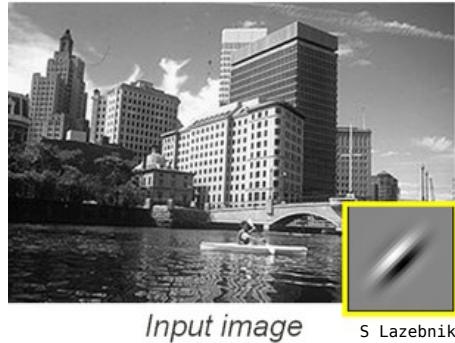
- Data preprocessing (& augmentation)
- Validation set
- Saving checkpoints during training
- Bigger models (10-100 million parameters)
- Normalisation, regularisation
- Lateral/top-down/skip connections
- Computational resources



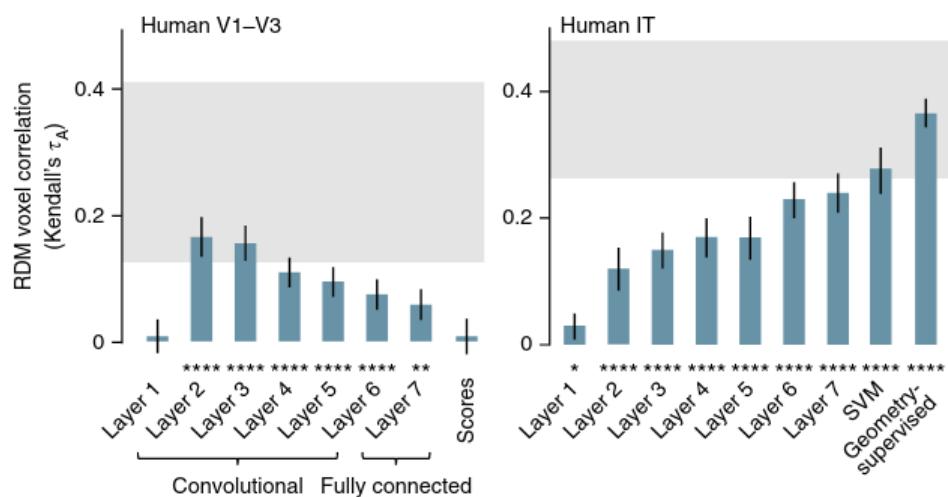
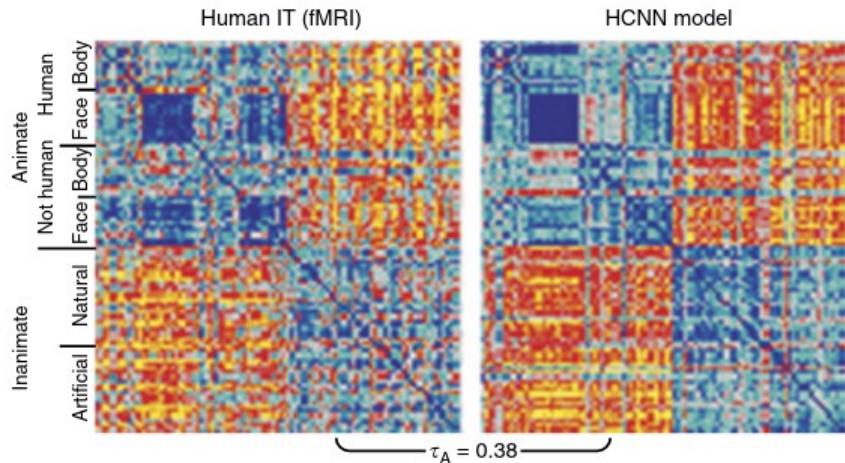
“Dog”



“Dog”



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]



Deep Neural Networks in Computational Neuroscience

FREE

Tim C. Kietzmann, Patrick McClure and Nikolaus Kriegeskorte

<https://doi.org/10>

Published online:

Perspective | Published: 28 October 2019

A deep learning framework for neuroscience

Blake A. Richards , Timothy P. Lillicrap, [...] Konrad P. Kording

Nature Neuroscienc

35k Accesses | 40

Abstract

The goal of computational neuroscience is to understand how the nervous system generates behavior. At the heart of this endeavor is the computational model of how sensory stimuli to neurons are processed. The range from simple to complex models has been dominated by “deep learning” approaches. In this “network” suggests that current DNNs are not able to capture the simplifications that are made in order to perform complex tasks such as perception and auditory scene analysis. In addition to the question of whether DNNs are predicting neural activity better than any other current model, there are other questions which are required to evaluate their performance. For example, black boxes, the lack of transparency, and the four directly mismatched between objective, and

Convolutional Neural Networks as a Model of the Visual System: Past, Present, and Future

Grace W. Lindsay

Perspective | Published: 16 November 2020

If deep learning is the answer, what is the question?

Andrew Saxe , Stephanie Nelli  & Christopher Summerfield 

Nature Reviews Neuroscience (2020) | [Cite this article](#)

4987 Accesses | 160 Altmetric | Metrics

Abstract

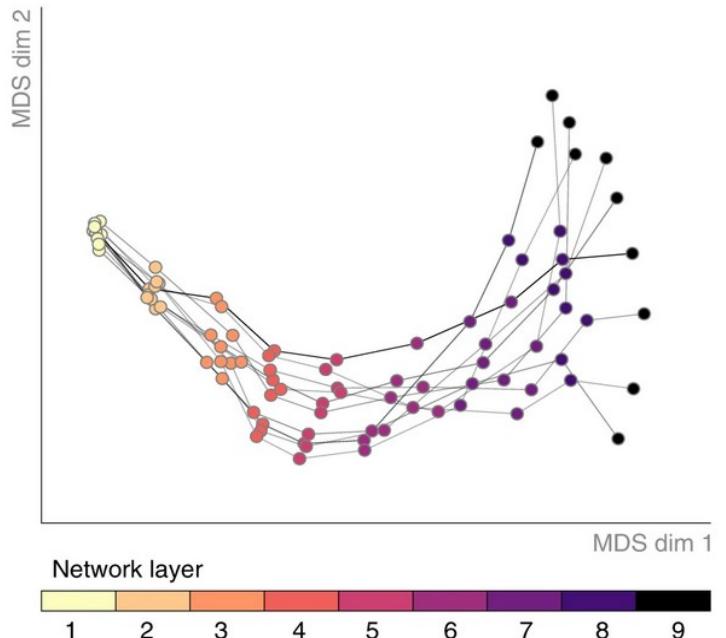
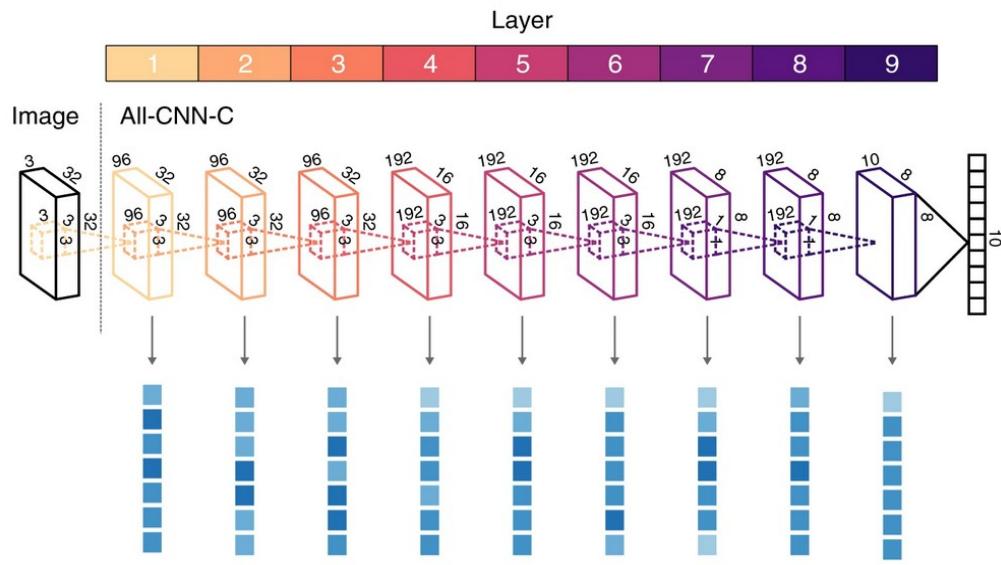
■ Convolutional neural networks (CNNs) have become the dominant tool for analyzing biological visual systems. Their success in computer vision has led to their widespread use in neuroscience. This Perspective highlights what, in the context of the visual system, CNNs can and cannot do, and how they can be used to build better models of the visual system. It also discusses the challenges of applying CNNs to the study of biological visual systems.

INTRODUCTION

Computational models serve several purposes in science. They can validate inferences drawn from experiments, predict new findings, drive theoretical developments, and provide a principled perspective on the field.

Abstract

Neuroscience research is undergoing a minor revolution. Recent advances in machine learning and artificial intelligence research have opened up new ways of thinking about



Models should be treated as participants

ecj35@cam.ac.uk

 emercjones

 github.com/emercjones/dnn-examples

